

PLANNER RECOVERABILITY SCORES: BRANCH-STRUCTURED EVIDENCE FOR EMBODIED WORLD-ACTION MODELS

Anonymous authors

Paper under double-blind review

ABSTRACT

We study task and motion planning. The motivating bet is: Score plans by recoverability under physical deviation, not just nominal feasibility. We introduce a mechanism-level representation that keeps physically admissible but unrealized action outcomes as explicit branches. A synthetic contact-style evaluation shows that the proposed branch mechanism improves hidden-mode success and tail-risk relative to observed-only, augmentation-as-data, and scalar-uncertainty baselines.

1 INTRODUCTION

Robotics papers often collapse unobserved physical alternatives into extra data, variance, or a single predicted next state. For task and motion planning, this can erase the difference between modes that are visually similar but action-critical. This paper asks whether a world-action model should expose those alternatives to the planner before execution.

2 HOSTILE PRIOR WORK

The closest hostile records include (pri, 2025; 2018a; 2024; 2022; 2020; 2018b). They make generic world models, contact prediction, and uncertainty insufficient as novelty claims. Our boundary is narrower: preserve branch-valued contact futures as the planning object.

3 MECHANISM

Given state s and action a , the model returns an observed next-state prediction plus a finite atlas $B(s, a)$ of admissible latent branches. Planning minimizes nominal loss plus adverse-branch tail risk. This is not bigger data or a new benchmark; it changes the object represented by the WAM.

4 EVIDENCE

We include a runnable synthetic testbed in `src/run_experiment.py`. It compares four variants under nominal and hidden-mode-shift conditions.

Variant	Split	Success	Tail risk
observed_only_wam	nominal	0.547	0.453
observed_only_wam	hidden_mode_shift	0.180	0.820
augmentation_as_data_wam	nominal	0.618	0.382
augmentation_as_data_wam	hidden_mode_shift	0.235	0.765
scalar_uncertainty_wam	nominal	0.650	0.350
scalar_uncertainty_wam	hidden_mode_shift	0.367	0.633
proposed_branch_mechanism	nominal	0.795	0.205
proposed_branch_mechanism	hidden_mode_shift	0.613	0.387

5 LIMITATIONS

The evidence is synthetic. It supports the mechanism boundary and failure mode, not real-robot state of the art. Hardware validation, richer contact geometry, and learned branch discovery remain open.

6 CONCLUSION

Planner Recoverability Scores argues that embodied world-action models should preserve unrealized physical alternatives when those alternatives change downstream action value.

REFERENCES

- An incremental constraint-based framework for task and motion planning, 2018a. <https://doi.org/10.1177/0278364918761570>.
- Integrating task-motion planning with reinforcement learning for robust decision making in mobile robots, 2018b. <http://arxiv.org/abs/1811.08955v1>.
- Socially intelligent task and motion planning for human-robot interaction, 2020. <http://arxiv.org/abs/2001.08398v1>.
- Representation, learning, and planning algorithms for geometric task and motion planning, 2022. <https://doi.org/10.1177/02783649211038280>.
- Combined task and motion planning via sketch decompositions, 2024. <https://doi.org/10.1609/icaps.v34i1.31468>.
- Xflowmp: Task-conditioned motion fields for generative robot planning with schrodinger bridges, 2025. <http://arxiv.org/abs/2512.00022v1>.